

COMP4300 Assignment 2

Jackson Kilrain-Mottram (u6940136)

May 2023

Just before you begin reading this, I wanted to note a few things about this assignment. I've had little to no time to actually work on this during the vast majority of the assignment duration (due to overlapping work and uni assessments that take up my time, hence not grounds for an extension). What has been implemented, evaluated, discussed and presented here is essentially just a week and a half worth of work. Is it optimal? No. Is it the best representation of my work? No. That being said, I think the ideas and techniques presented are well thought through. Anyway, cheers for your consideration, hopefully this isn't too painful to read.

1 Question 1

1.1 Performance

In order to properly quantify the performance impacts that the various configurations of OpenMP directives used in each case have, we will utilise a Linux tool called **perf**. This captures event statistics for the compute stack, from the OS level, to kernel and hardware. We are particularly interested in hardware, at the L1 data cache misses and LLC load misses. To give some context to these quantities we will examine, let's first mention that the program has roughly 5.4×10^{10} instruction total to be executed for 100 iterations of a 1000×1000 grid.

```
1 void omp1dUpdateBoundary(...) {
2     // ... snip ...
3     #pragma omp for private(j)
4     for (j = 1; j < N + 1; j++) { ... }
5     #pragma omp for private(i)
6     for (i = 0; i < M + 2; i++) { ... }
7 }
8
9 void omp1dUpdateAdvectField(...) {
10    // ... snip ...
11    #pragma omp for private(i)
12    for (i = 0; i < N; i++)
13        for (j = 0; j < M; j++)
14            V(u, i, j) = ...;
15 }
16
17 void omp1dCopyField(...) {
18    // ... snip ...
19    #pragma omp for private(i)
20    for (i = 0; i < N; i++)
21        for (j = 0; j < M; j++)
22            V(u, i, j) = V(v, i, j);
23 }
24
25
26 void omp1dAvect(...) {
27    // ... snip ...
28    for (size_t r = 0; r < reps; r++) {
29        #pragma omp parallel shared(u,ldu,v,ldv)
30        {
31            omp1dUpdateBoundary(u, ldu);
32            omp1dUpdateAdvectField(&V(u, 1, 1), ldu, &V(v, 1, 1), ldv);
33            omp1dCopyField(&V(v, 1, 1), ldv, &V(u, 1, 1), ldu);
34        }
35    }
36 }
```

Utilising the performance counters from **perf stat -d**, we see that on average **L1-dcache-load-misses** is approximately 0.72% and **LLC-load-misses** fluctuate between 55 – 62.08%. Quantitatively, these equate to 5.6×10^7 in L1 dcache misses and 1.4×10^6 in LLC misses. This implies that there are few misses in the cache for loads, as the locality and alignment of data shared between threads is well situated. The overhead seen here can be attributed to the edge cases between the outside (starting and ending point) of the contiguous part of the advection field accessed for a thread. These are less ideal in their alignment and block saturation, leading to the minimum overhead seen in the L1 data cache load misses.

1.2 Parallel Region Entry/Exits

```
1 void omp1dUpdateBoundary(...) {
2     // ... snip ...
3     #pragma omp for private(j)
4     for (j = 1; j < N + 1; j++) { ... }
5     #pragma omp for private(i)
6     for (i = 0; i < M + 2; i++) { ... }
7 }
8
9 void omp1dUpdateAdvectField(...) {
10    // ... snip ...
11    for (i = 0; i < N; i++)
12        #pragma omp for private(j)
13        for (j = 0; j < M; j++)
14            V(u, i, j) = ...;
15 }
16
17 void omp1dCopyField(...) {
18    // ... snip ...
19    for (i = 0; i < N; i++)
20        #pragma omp for private(j)
21        for (j = 0; j < M; j++)
22            V(u, i, j) = V(v, i, j);
23 }
24
25 void omp1dAvect(...) {
26    // ... snip ...
27    for (size_t r = 0; r < reps; r++) {
28        #pragma omp parallel shared(u,ldu,v,ldv)
29        {
30            omp1dUpdateBoundary(u, ldu);
31            omp1dUpdateAdvectField(&V(u, 1, 1), ldu, &V(v, 1, 1), ldv);
32            omp1dCopyField(&V(v, 1, 1), ldv, &V(u, 1, 1), ldu);
33        }
34    }
35 }
36 }
```

1.3 Cache Misses Coherent Reads

TODO

1.4 Cache Misses Coherent Writes

```
1 void omp1dUpdateBoundary(...) {
2     // ... snip ...
3     #pragma omp for private(j)
4     for (j = 1; j < N + 1; j++) { ... }
5     #pragma omp for private(i)
6     for (i = 0; i < M + 2; i++) { ... }
7 }
8
9 void omp1dUpdateAdvectField(...) {
10    // ... snip ...
11    #pragma omp for private(i) schedule(dynamic)
12    for (j = 0; j < N; j++)
13        for (i = 0; i < M; i++)
14            V(u, i, j) = ...;
15 }
16
17 void omp1dCopyField(...) {
18    // ... snip ...
19    #pragma omp for private(i) schedule(dynamic)
20    for (j = 0; j < N; j++)
21        for (i = 0; i < M; i++)
22            V(u, i, j) = V(v, i, j);
23 }
24
25 void omp1dAvect(...) {
26    // ... snip ...
27    for (size_t r = 0; r < reps; r++) {
28        #pragma omp parallel shared(u,ldu,v,ldv)
29        {
30            omp1dUpdateBoundary(u, ldu);
31            omp1dUpdateAdvectField(&V(u, 1, 1), ldu, &V(v, 1, 1), ldv);
32            omp1dCopyField(&V(v, 1, 1), ldv, &V(u, 1, 1), ldu);
33        }
34    }
35 }
36 }
```

```

34 }
35 }
36 }

```

Using `OPENMP_NUM_THREADS=48 perf stat -d ./testAdvect 1000 1000 100` we can keep track of performance counters of hardware behaviour while the advection solution is executing. More specifically, paying attention to the `L1-dcache-load-misses` and `LLC-load-misses`. From the execution profile, we note that `L1-dcache-load-misses` sits at $6.80 - 14.60\%$ miss rate consistently across multiple runs. For `LLC-load-misses`, this tends to fluctuate heavily between 75.73% overall. From this we can see that cache misses occur frequently for loads that are localised to blocks that overlap between processors. This causes frequent invalidation due to coherence reads between the processors to ensure up to date cache block data. Qualitatively, there is approximately 5.9×10^8 total L1 dcache misses, compared to the 5.6×10^7 for the optimise variation in case 1. Similarly the LLC misses are an order of magnitude greater in quantity, at 1.35×10^7 on average compared to 1.4×10^6 .

1.5 Performance Analysis

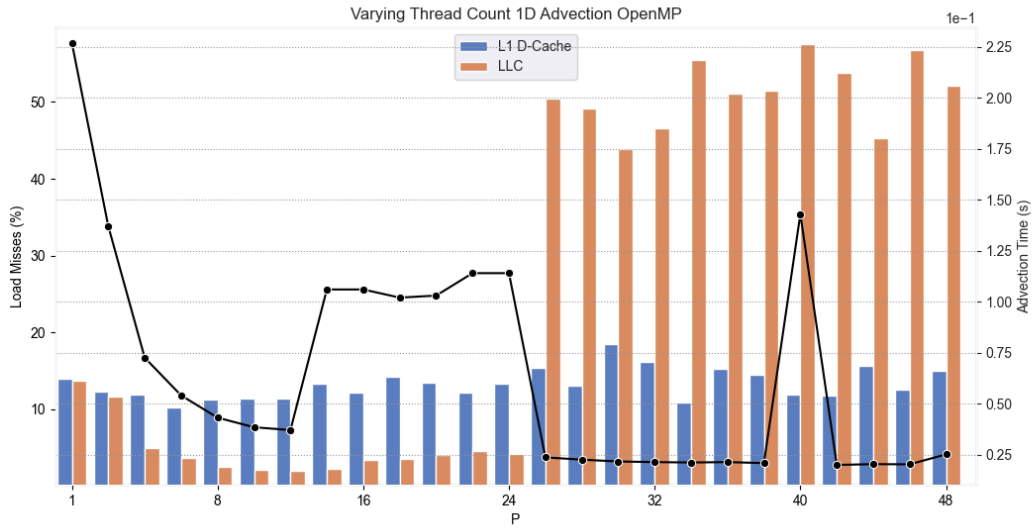


Figure 1: 1D advection time and cache misses (L1 data and LLC) for varying thread counts. Line graph uses right y-axis, bar plots use left y-axis.

TODO: Write up analysis here

2 Question 2

In order to formulate a model, we need to establish a few points on the hardware properties and problem decomposition. Firstly, using `cat /proc/cpuinfo | grep cache_alignment` we can determine the cache block size for the Xeon 8274 CPUs being utilised. Using this on the Gadi nodes, we see that there is 64 bytes in a cache block. To set the scene, an $M \times N$ grid, requires $\lceil \frac{M \cdot N}{64} \rceil$ cache blocks to fully contain it, these are distributed across the contiguous grid (figure 2).




Figure 2: Cache block distribution over an example 2D advection grid.

In order to determine cache coherence behaviour on read and write operations for the solution, we need to first define a function $f : \mathbb{Z}^2 \rightarrow \mathbb{Z}$ that maps a 2D grid index (i, j) to the associated cache block it resides in.

$$f(i, j) \stackrel{\text{def}}{=} \left\lfloor \frac{j + (i \cdot N)}{64} \right\rfloor$$

2.1 Field Update

For the `omp1dUpdateAdvectField` method, there are two sets of operations to consider for coherent-write cache misses and coherent-read cache misses.

2.1.1 Write Misses

Each row starts at a given cache block and ends at the same or another cache block. Using this we can determine how many cache lines (and which they are) are involved. For each write, we need to invalidate these cache lines on each of the other processors. Given that each of the i indices are in parallel, we need to find the maximal cost line update across all processors. Therefore the total can be expressed as follows

$$M_L = \left\lfloor \frac{M}{P} \right\rfloor$$

$$U_w = t_{w,W} \cdot \max_{p \in [0..P]} \left(\max \left\{ 1, \sum_{i=1}^{M_L} [f((p \cdot M_L) + i, N) - f((p \cdot M_L) + i, 1)] \right\} \right)$$

2.1.2 Read Misses




Figure 3: First and last cache block encounters over a subset of the example advection grid.

Each row element has a convolution applied over the previous, current and next row elements. The rows start at a given cache block and end at the same or different cache block. For each transition between cache blocks (as highlighted in figure 3), we will see a coherent read miss. Given we load these rows in chunks for each thread, we consider misses for each row, summed together.

$$U_r = t_{w,R} \cdot \max_{p \in [0..P]} \left(\max \left\{ 1, \sum_{i=1}^{M_L} \left(\sum_{l=-1}^1 [f((p \cdot M_L) + i + l, N) - f((p \cdot M_L) + i + l, 1)] \right) \right\} \right)$$

2.2 Copy Field

When copying the field from v to u to allow for the next iteration or finalisation to occur, we have both read and write operations occurring simultaneously. Note that these are orthogonal to each other, as the read operations are performed over v and the write operations over u . We can abstract the common max into a single formulation applied to both the read and write coefficients as they equate equivalently.

$$C_{\max} = \max_{p \in [0..P]} \left(\max \left\{ 1, \sum_{i=1}^{M_L} [f((p \cdot M_L) + i, N) - f((p \cdot M_L) + i, 1)] \right\} \right)$$

$$C_w = t_{w,W} \cdot C_{\max}$$

$$C_r = t_{w,R} \cdot C_{\max}$$

2.3 Full Model

Combining the above operational breakdowns, we can achieve the full model over r iterations with an entry/exit cost for $U_{r,w}$ and $C_{r,w}$, ignoring the boundary update (as stipulated by the question):

$$t = r \cdot (2t_s + U_r + U_w + C_r + C_w)$$

2.4 Coefficient Evaluation

Considering the measurement for t_s , we can use two times, one that captures the startup overhead and the second that captures the teardown overhead. For a given parallelised loop, this equates to computing the time to the first loop iteration and the time from the last loop iteration. The following pseudo code illustrates this.

```

1 time startup_t1, startup_t2;
2 time teardown_t1, teardown_t2;
3 startup_t2 = getCurrentTime();
4 #pragma omp parallel for
5 for (int i = 0; i < N; i++) {
6     if (i == 0) {
7         startup_t2 = getCurrentTime();
8     } else if (i == N) {
9         teardown_t1 = getCurrentTime();
10    } else {
11        // do stuff
12    }
13 }
14 teardown_t2 = getCurrentTime();

```

Performing this within the advection solver, we get an average value of 6.74×10^{-4} s for the total startup + teardown overhead. Comparing the to value measured in assignment 1 ($2.951666666667 \times 10^{-4}$), it is indeed close to the cost of startup for an MPI communication. More specifically, it is close to the two way communication startup cost. If we consider only the startup, it is comparatively close to the one way cost at 1.6803624×10^{-5} against 1.47583×10^{-5} (MPI).

We can easily inspect the values of $U_{r,w}$ and $C_{r,w}$ by splitting up each inner most loop into two distinct operations. The first performs reads, and the latter performing writes, where the measured times are written into thread-local array entries. We can then average these to find the coefficients for each of the operation overheads between the iterations that require coherence updates and the ones that don't (reads for instance will require an update on the first iteration indexing a new cache line and writes on every iteration).

	Read ($t_{w,R}$)	Write ($t_{w,W}$)
C	6.26285×10^{-6}	2.220465×10^{-5}
U	2.125284×10^{-5}	5.465016×10^{-5}
Average	1.3757845×10^{-5}	3.8427405×10^{-5}

Table 1: Measurements of t_w without timer overheads for update U and copy C parallelised operations with 24 threads and $p = 12$. Note that the startup and teardown time t_s is not included in these measurements as they are from within the iterations.

Compared to the value of t_w from assignment 1 ($6.469449853897094 \times 10^{-10}$), the value calculated here is three orders of magnitude greater. Given that the behaviour of cache coherency is limited to hardware effects on the same node, this is unusually. Especially considering the communication cost for a single double for the MPI implementation is over a network. It is possible that these measurements are skewed by hardware behaviour or low-level OS behaviour such as context switches and other of the like.

3 Question 3




Figure 4: 2D advection time and cache misses (L1 data and LLC) for varying aspect ratios. The first group (left of left red line), are square ratios, the second group (in between red lines) are ratios for 48 threads and the last right group (right of right red line) are for 24 threads. Line graph uses right y-axis, bar plots use left y-axis.

TODO: Write up analysis here

4 Question 4

4.1 Solution Description

For convolution problems that compute time step evolutions, an approach for very large scale matrices is to utilise Fourier transforms (?). In assignment 1, this was also presented and translated into an MPI implementation based on the original algorithm in the paper by ?. Here, the same algorithm has been implemented, utilising OpenMP as the parallelisation medium. There were several issues with the initial implementation in assignment 1 that have been fixed; missing final rotation of result matrix, incorrect forward DFT parameters and a few other minor points.

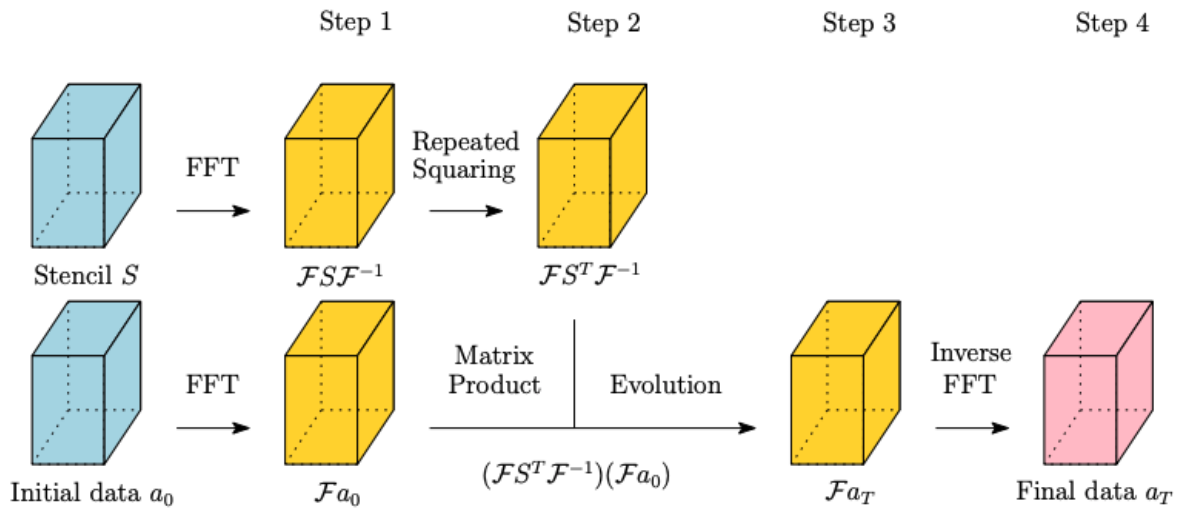


Figure 5: Outline of applying a FFT convolution kernel to an iterated stencil computation tile. Image sourced from ?.

As an overview, this approach applies time step evolution as repeated squaring to the FFT transformed core convolution kernel (figure 5). Given the sparse nature of the matrix, this vastly simplifies the total quantity of matrix-matrix operations required to compute the full evolution. Computing a single matrix-matrix product with the FFT evolved kernel and the FFT evolved data matrix, the full evolution can be computed with an overall smaller computational overhead. Lastly, the inverse FFT is computed and the result rotated $T \bmod M$ (where T is the time steps and M is

the matrix width).

We can utilise a convolution kernel derived from the Lax-Wendroff method using the coefficient functions of the Courant number used in the definition as described in ?. Essentially, we end up with the following kernel (see the assignment 1 report for further details and motivation on this):

$$\begin{aligned}
c_y &= a \frac{\Delta t}{\Delta y} \\
b_{-1}^y &= \frac{c_y}{2}(1 + c_y) \\
b_0^y &= 1 - c_y^2 \\
b_1^y &= -\frac{c_y}{2}(1 - c_y) \\
\mathcal{M} &= \begin{bmatrix} b_{-1}^x \cdot b_{-1}^y & b_{-1}^x \cdot b_0^y & b_{-1}^x \cdot b_1^y \\ b_0^x \cdot b_{-1}^y & b_0^x \cdot b_0^y & b_0^x \cdot b_1^y \\ b_1^x \cdot b_{-1}^y & b_1^x \cdot b_0^y & b_1^x \cdot b_1^y \end{bmatrix}
\end{aligned}$$

$$\forall i \in [0, m], j \in [0, n] \quad u_{i,j}^{n+1} = u_{i,j}^n * \mathcal{M}$$

NOTE: In order to compile and use this, provide the `FFT_CONV_KERNEL=1` flag alongside the make command, e.g. `module load fftw && make clean FFT_CONV_KERNEL=1 all`.

4.2 Performance Evaluation

As mentioned in the previous sub-section, this FFT method is designed to work with very large evolutionary convolution problems. In order to properly illustrate this, let us consider a matrix of dimensions $10,000 \times 10,000$ over 100 time steps. For both the optimised evaluation and the baseline, we consider 24 threads using a 6×4 decomposition on a single NUMA node to ensure minimal overhead in cache coherence behaviour (as explored in section 2 and 3).

For the baseline we see that the advection time is 1.19×10^2 s, which is substantially larger than any of the previous results in this report. Utilising the FFT optimisation, this drops to 2.57×10^1 s, equating to a change of -78.4% overall. One could argue that this is still a long time to compute the advection of 100 time steps, however considering that each thread is responsible for a chunk of $\sim 1667 \times 2500$, this is a huge advantage that indicates the standard approach is heavily compute bound in quantity of arithmetic operations.

5 Question 5

NOTE: Compare the following

- compare the speedups for our parallel baseline code against 1 GPU core
- also compare the speedups for our parallel baseline code against 1 host core

TODO: Measure advection time, achieved warp occupancy and throughput:

```
nv-nsight-cu-cli --metrics achieved_occupancy,gld_throughput ./testAdvect -g Gx,Gx -b Bx,By M N r
```

6 Question 6

TODO:

Uses cuda streams to concurrently interleave warp execution of both the boundary and inner field.

Does not create dedicated stream for only the corner since it doesn't fill enough of a single warp, which degrades performance further than just dealing with the minor amount of additional warp divergence/converged from additional conditional logic on corners in boundary updates.

7 Question 7

Between the three models of MPI, OpenMP and CUDA, design wise CUDA provides the greatest diversity for design of highly parallel tasks. More specifically, with regards to problems that have thread independent discretisation models, such as matrix evolutions (Heat equation, advection, etc). However, when it comes to approaches that require interaction between threads (not necessarily synchronisation), it introduces complexity that is otherwise more manageable in the MPI model. As for OpenMP, it yields similar benefits that CUDA does, however in much more significant context given the restrictions on compiler driven threading models as opposed to the more explicit nature of CUDA.

Implementation wise, OpenMP has the most approachable interfaces for both new and adaptive implementation. This extends to understanding the behaviour of how execution, compaction and grouping function relative to the behaviour of the pragma directives. That being said, it has a high-end implementation overhead when it comes to optimising it for fine-grained control of execution behaviour due to the high-level of abstraction. MPI has provides a well rounded SPI from simplistic interaction models for control over exchange behaviour and optimisations to a per-thread level to the point of fine-grained memory behaviour.

Over CUDA, it is substantially easier to control complex overlapping behaviour where the implementation does not follow SIMD/SIMT structuring well (diversity in condition evaluation for instance). This allows for reduced algorithmic complexity when handling these kinds of diverse behaviour on a thread-to-thread basis. Lastly, for CUDA, implementation has a high-entry barrier in terms of understanding how to appropriately structure your implementation and understand the resource decomposition required for appropriate allocation and usage. This requires more work to structure the implementation differently, especially when considering how conditional overheads (warp divergence/convergence) are easy to introduce with easy-to-overlook implementation aspects. Substantially more effort and understanding is required for the CUDA model, however, even with this overhead the tool set it provides for programmability is far and above the others allowing full control over the behaviour and program structure across threads when compared to OpenMP.

During development, in any of the three paradigms/models, the debugging of the core algorithm with regards to structure, is essentially the same. This is specific to the arithmetic operations and memory operations. However, considering the halo exchanges and the different approaches for managing these, it differs quite a bit ween the communication model versus the shared memory model. For MPI, exchange debugging relies on logging tagged entries communicated between threads for message issues. With contention issues much of the debugging hinges on manual checking and verification of parameters for communication handlers.

On the other hand, in the land of shared memory models, OpenMP is straightforward to debug for most cases. In the most difficult of cases (that I had encountered), it comes down to data ownership between threads and the global program space. Missing a marker for shared or private variables can lead to undefined behaviour that is hard to track down, however, due to the simplicity of compiler directives for simple cases (such as advection solution), these tend to be easy to find. Compare that to CUDA, which has improvement with regards to scoping as mentioned previously, as runtime errors occur for segmentation violations and improper host/device pointer usage. These become much clearer, however, sourcing the problem behind them can be arduous for complex logic. With proper alteration of kernel execution and isolation, it does however provide an edge over OpenMP with the clarity of debugging techniques.

References